HTTP Caching Proxy
Engineering Robust Server Software
Homework 2

For this assignment you will be writing an http proxy – a server whose job it is to forward requests to the origin server on behalf of the client. Your proxy will cache responses, and, when appropriate, respond with the cached copy of a resource rather than re-fetching it.

While the HTTP specification is quite large (and includes many complex features), you should make an http proxy which functions with GET, POST, and CONNECT. (You MAY handle any of the other request methods if you want). Specifically, a user should be able to configure their browser to use your proxy, and browse typical webpages (e.g., perform a Google Search, view the results, etc). Note that many webpages only do HTTPS, so your browser will use CONNECT to communicate with them (and you won't see the actual GET requests).

Your proxy MUST cache responses (when they are 200-OK) to GET requests. You should follow the rules of expiration time and/or re-validation in determining if your proxy can serve a request from its local cache (versus re-fetching from the origin server). Other cache management policies (e.g., replacement policy) are up to you.

Your proxy MUST be able to handle multiple concurrent requests effectively and SHOULD use multiple threads as part of your strategy to do so. The remainder of the design of handling multiple requests is up to you, but your cache MUST be shared between all connections (and properly synchronized).

Your proxy MUST produce a log (in `/var/log/erss/proxy.log`) which contains information about each request. To keep the log understandable, your proxy will assign each request a unique identifier when it prints the first log message for that request.

In each of the following format descriptions `typewriter text` is literal, and *italics indicate a variable*.

- Upon receiving a new request, your proxy should assign it a unique id (*ID*), and print the ID, time received (*TIME*), IP address the request was received from (*IPFROM*) and the HTTP request line (*REQUEST*) of the request in the following format:
  ID`: "`*REQUEST*`" from `*IPFROM* `@ `*TIME*

- If the request is a GET request, your proxy should check its cache, and print one of the following:
  ID`: not in cache`
  ID`: in cache, but expired at `*EXPIREDTIME*
  ID`: in cache, requires validation`
  ID`: in cache, valid`

- If your proxy needs to contact the origin server about the request, it should print the request it makes to the origin server:

  ID: `Requesting "`*REQUEST*`" from` *SERVER*

  Later, when it receives the response from the origin server, it should print:

  ID: `Received "`*RESPONSE*`" from` SERVER

  Here, *REQUEST* and *RESPONSE* are the request line and response line (first line in the message), and *SERVER* is the server name.

- If your proxy receives a 200-OK in response to a GET request, it should print one of the following:

  ID: `not cacheable because` *REASON*

  ID: `cached, expires at` *EXPIRES*

  ID: `cached, but requires re-validation`

- Whenever your proxy responds to the client, it should log:

  ID: `Responding "`*RESPONSE*`"`

  Where response is the response line of the reply. Note that you should do this if you reply with an error (e.g. if you receive a malformed request).

- When your proxy is handling a tunnel as a result of 200-OK, it should log (in addition to all other normal logging) when the tunnel closes with

  ID: `Tunnel closed`

- Your proxy MAY include any other log messages of your choice, as long as they adhere to the following formats:

  ID: `NOTE` *MESSAGE*

  ID: `WARNING` *MESSAGE*

  ID: `ERROR` *MESSAGE*

  Where MESSAGE is text of your choice (not containing a new-line). You may use (`no-id`) as the ID if there is no id for one of these.

<mark>All times should be printed in UTC, with a format given by `asctime`.</mark> For example, your log might say:

```
104: "GET www.bbc.co.uk/ HTTP/1.1" from 1.2.3.4 @ Sun Jan 1 22:58:17
2017
104: not in cache
105: "GET www.duke.edu/foo/bar HTTP/1.1" from 11.12.42.40 @ Sun Jan 1
22:58:17 201
104: Requesting "GET www.bbc.co.uk/ HTTP/1.1" from www.bbc.co.uk
105: in cache, valid
105: Responding "HTTP/1.1 200 OK"
104: Received "HTTP/1.1 200 OK" from www.bbc.co.uk
104: NOTE Cache-Control: must-revalidate
104: NOTE ETag: W/"33bc8-F9Kn1zgYX0cHOaRFsmZORA"
104: cached, but requires re-validation
```

```
(no-id): NOTE evicted www.foo.bar.com/boring.txt from cache
104: Responding HTTP/1.1 200 OK
```

In addition to the functional requirements specified above, there are some other requirements for this assignment:

- You MUST implement your proxy in C++. Please make good use of OO design, RAII, exceptions, and other C++-concepts. Do NOT write C and call it C++.

- Your proxy MUST be robust to external failures. If it contacts the destination webserver and receives an error response, it must handle it gracefully. If it contacts the destination webserver and receives a corrupted response, it MUST reply with a 502 error code. If the proxy receives a malformed request, it MUST reply with a 400 error code.

- You SHOULD think carefully about the exception guarantees you make, and how you handle problems. You SHOULD document these in your code.

- You MUST provide a docker-compose.yml file and a Dockerfile (or more if appropriate) which allow your proxy to be run with sudo docker-compose up. The Docker/Docker-Compose setup that you create MUST do the following:

  - Connect the host computer's port 12345 to your proxy.
  - Mount a directory called logs (in the same directory as the docker-compose.yml file) to /var/log/erss in the container running your proxy. Given that your proxy will write its logs to /var/log/erss/proxy.log inside the container, this means you (and your TAs) should be able to find proxy.log in the logs directory on the host.

- You MUST provide a set of testcases which demonstrate the functionality of your proxy – both in the common case, and in the case where it handles errors or unusual situations. You may find netcat incredibly useful here as it will let you send malformed requests or responses. You may also find wget useful as it will generate standard web requests (and can print the request, giving you a starting point for a malformed request). We encourage you to provide an automated/self-contained setup to demonstrate these testcases. You might consider making use of something like docker swarm to show how your proxy handles many concurrent requests.